

## Problem 1. Implementing a Neural Network Classifier

In this homework, you will implement a neural network and the backpropagation algorithm and stochastic gradient descent with mini-batches. The goal is to implement this yourself with only basic linear algebra packages such as *numpy*, rather than a deep learning library. We will use MNIST, a common classification data set consisting of 28x28 pixel grayscale images of handwritten digits (0-9), with 60,000 training examples and 10,000 test examples. The task is to predict the digit from the image, using a *softmax* output layer.



Figure 1: MNIST example images

### Requirements:

- Use tanh or relu hidden units. Logistic (also called sigmoid) units don't work as well in the hidden layers.
- Use mini-batch gradient descent (e.g. mini-batches of size 100). You can use momentum or any other tricks you wish to speed up training.
- Train a network with at least two hidden layers, and try to tune the hyperparameters (e.g. learning rate) get good performance on the test data. Plot both the training and test *loss* (the NLL) throughout training (at every epoch), and write a paragraph describing the network, training algorithm, and any hyperparameter choices. Make sure to report the *accuracy* (also called the *01-loss* — to get a class prediction from a softmax layer, we simply take the *argmax* of the probabilities).
- Next, try to construct a network that *overfits* to the training data, so that the test loss increases while the training loss continues to decline. (Hint: this should be easier with a network of a single hidden layer.)

### Tips:

- For a minimalist implementation of a neural network, see:  
<https://iamtrask.github.io/2015/07/12/basic-python-network/>
- To quickly download the MNIST data set, see:  
<https://github.com/hsjeong5/MNIST-for-Numpy>

- (c) A single pass through the training data is called an *epoch*. You may have to do 50-100 epochs to see overfitting. If your computer is slow, go ahead and use a fraction of the data set, e.g. 6,000 examples.
- (d) We typically compute the validation loss at the end of every training epoch to see if we are overfitting. For the training data set, we usually keep a running average of the loss while performing the mini-batch training updates to save on computation.